

Rahul Prakash



(+91) 9688107002



rahulmurugan96@gmail.com



/rahulprakash

*Building Data Lakes @ Fivetran | 7+ years | Ex-OCI and Zoho | Backend Development | Java |
Distributed Systems | Microservices | SaaS | AWS | GCP | Azure*

EXPERIENCE

Fivetran India Pvt Ltd Staff Software Engineer

Bangalore, India
September 2025 – present

Iceberg Catalog

- Drive improvements in monitoring, observability, and reliability for the Apache Polaris open-source catalog.

Senior Software Engineer

March 2023 – August 2025

Datalake Destinations:

- Led the successful onboarding of diverse new Datalake destinations (**MSK, Aiven Kafka, Microsoft OneLake**), significantly expanding Fivetran's platform capabilities.
- Integrated critical metastore and catalog solutions (**BigQuery Metastore, Snowflake Open Catalog, Iceberg Hadoop Catalogs**) to enhance data interoperability and management.
- Achieved substantial performance enhancements in Data Lake operations by optimizing DuckDB queries and parallelizing data merge workflows. **Improved upload and download latency by 50% and reduced data load times by 70%.**
- Developed and implemented a robust framework for automated Delta Lake table maintenance, including purging old data files, optimizing customer storage costs and data hygiene.
- Designed and engineered comprehensive support for the Iceberg table format within Azure ADLS storage, extending core data lake functionalities.

Fivetran In-House Catalog:

- Architected and built a new in-house catalog leveraging **open-source technologies** to serve as a single source of truth for data metadata.
- Developed a robust framework for customer migration and onboarding to the new Fivetran catalog, ensuring seamless transition.
- Integrated core Fivetran internal services with the in-house catalog to enable seamless interoperability and consistent data discoverability across the platform.

Oracle India Pvt Ltd Software Development Engineer II

Nov 2021 – March 2023
Bangalore, India

Oracle Cloud Infrastructure – Data Science

- Designed and implemented **ETag** for resource-specific APIs, enabling **conditional requests and concurrency control**.
- Optimized document processing via preview generation and content API streaming, **reducing memory utilization by 75%.**
- Automated end-to-end Disaster Recovery setup, cutting setup time and eliminating manual processes.
- Built **dynamic password rotation** system for database users, enhancing critical security posture.
- Developed essential **utility tools** for automated password rotation and secure secrets updates to the vault service, streamlining security operations and eliminating manual overhead.
- Delivered critical hotfixes and monthly deployments, maintaining system stability and SLA compliance.
- Resolved complex API issues and managed library updates to ensure reliability and compliance.

Creator Onpremise – Multi tenant low code application building platform.

- Worked on building OnPremise version of Zoho Creator cloud product from scratch.
- Developed a framework to upgrade from old version of product to new.
- Integrated **Authentication service** for unified user experience.
- Designed and implemented pricing module of OnPremise version in order to provide an unified user experience in Cloud and OnPremise.
- Designed and Developed Onpremise version build with the help of **Apache ANT**.
- Worked on various performance tasks like reducing the application loading latency from **30s to 30ms**, reducing the component level access **latency from 18s to 900ms and reducing latency by 67% and memory utilization by 41% in user key cache storage**.
- Have actively involved in code review process and bug fixing.

On Premise – Multiple Application Server

- Designed and developed a POC model for Multiple Application server.
- Integrated **Redis** and Application was made stateless by storing static variables in property file and moved in-memory cache to Redis.
- Worked in balancing load to the Multiple Application Servers using **Nginx** and deployment of load balancers and application servers.

EDUCATION

Madras Institute of Technology

Bachelor of Engineering in Computer Science and Engineering

2014 – 2018
Chennai, India

Coursework: Data Structures, Analysis of Algorithms, Database Management Systems, Object Oriented Analysis and Design, Operating Systems, Computer Networks

TECHNICAL SKILLS

Programming Languages: Java, C++, Python

SQL: MySQL, PostgreSQL, ATP, DuckDB

NoSQL: Cassandra, Redis, DynamoDB

Data Platforms & Lakehouse: Delta Lake, Iceberg, OneLake, ADLS, S3, GCS, Hadoop Catalog, Snowflake Open Catalog, Apache Polaris, Databricks, BigQuery, Snowflake, DuckDB

Cloud Platforms: AWS, Azure, GCP, OCI

Message Queues: Apache Kafka

Frameworks: NGINX, Struts2, REST APIs, Apache ANT, Servlet, JSP, Spring Boot, Mockito, JUnit, Liquibase, gRPC

AWARDS AND RECOGNITION

- 2nd Place in Hackathon for implementing a connector for Snowflake Iceberg table format, which was later integrated as a Data Lake destination.
- Received Spot Recognition Award twice for outstanding contributions to the OneLake and Fivetran Catalog projects.
- Participated and won 3rd place in Oracle hackathon for implementing ORM analyser and release timeline predictor.
- Received annual technical excellence award in Zoho Corporation.
- Ranked top 10 in inter-college coding challenge participated by 1000+ students.