

# RAHUL PRAKASH



(+91) 9688107002



rahulmurugan96@gmail.com



/rahulprakash

## PROFESSIONAL SUMMARY

Staff Software Engineer with 8 years of experience in Backend Development, Distributed Systems, and Microservices, currently building Data Lakes at Fivetran. Proven expertise in architecting and developing scalable distributed systems with strong proficiency in Java, cloud platforms (AWS, GCP, Azure), and data lake technologies including Apache Iceberg and Delta Lake. Experienced in leading complex technical projects, optimizing system performance, and contributing to open-source projects. Adept at leveraging modern AI-powered development tools (GitHub Copilot, Claude, Gemini) to enhance productivity and code quality.

## TECHNICAL SKILLS

**Languages:** Java, Python, SQL, JavaScript, Shell Script

**Cloud Platforms:** AWS, Google Cloud Platform (GCP), Microsoft Azure, Oracle Cloud Infrastructure (OCI)

**Data & Storage:** Apache Iceberg, Delta Lake, Apache Kafka, DuckDB, PostgreSQL, MySQL, Cassandra, Redis, Azure Data Lake Storage (ADLS), Amazon S3

**Frameworks & Tools:** Spring Boot, Dropwizard, Quarkus, Docker, Kubernetes, Nginx, Git, Maven, Gradle

**Monitoring & Observability:** OpenTelemetry, Prometheus, Grafana, Opsgenie

**Architecture & Design:** Distributed Systems, Microservices, Data Lakes, RESTful APIs, System Design, Performance Optimization

**AI-Powered Development:** GitHub Copilot, Claude AI, Gemini, VS Code

## PROFESSIONAL EXPERIENCE

### STAFF SOFTWARE ENGINEER

September 2025 - Present

Fivetran India Pvt Ltd, Bangalore, India

#### Iceberg Catalog

- Drive improvements in monitoring, observability, and reliability for the Apache Polaris open-source catalog
- Lead architectural decisions and technical direction for catalog services supporting enterprise data lakes
- Contribute to open-source Apache Polaris project using Dropwizard and Quarkus frameworks

### SENIOR SOFTWARE ENGINEER

March 2023 - August 2025

Fivetran India Pvt Ltd, Bangalore, India

#### Datalake Destinations

- Led the successful onboarding of diverse new Datalake destinations (MSK, Aiven Kafka, Microsoft OneLake), significantly expanding Fivetran's platform capabilities
- Integrated critical metastore and catalog solutions (BigQuery Metastore, Snowflake Open Catalog, Iceberg Hadoop Catalogs) to enhance data interoperability and management
- Achieved substantial performance enhancements in Data Lake operations by optimizing DuckDB queries and parallelizing data merge workflows; **improved latency by 50% and reduced data load times by 70%.**
- Developed and implemented a robust framework for automated Delta Lake table maintenance, optimizing customer storage costs and data hygiene
- Designed and engineered comprehensive support for the Iceberg table format within Azure ADLS storage.

#### Fivetran In-House Catalog

- **Architected and built a new in-house catalog** leveraging open-source technologies to serve as a single source of truth for data metadata
- Developed a robust framework for customer migration and onboarding to the new Fivetran catalog, ensuring seamless transition
- Integrated core Fivetran internal services with the in-house catalog to enable seamless interoperability across the platform

# SOFTWARE DEVELOPMENT ENGINEER II

November 2021 - March 2023

Oracle India Pvt Ltd, Bangalore, India

## Oracle Cloud Infrastructure - Data Science

- Designed and implemented ETag for resource-specific APIs, enabling conditional requests and concurrency control
- Optimized document processing via preview generation and content API streaming, **reducing memory utilization by 75%**
- **Automated end-to-end Disaster Recovery setup**, cutting setup time and eliminating manual processes
- Built dynamic password rotation system for database users, enhancing critical security posture
- Developed essential utility tools for automated password rotation and secure secrets updates to the vault service
- Delivered critical hotfixes and monthly deployments, maintaining system stability and SLA compliance

## MEMBER TECHNICAL STAFF

May 2018 - October 2021

Zoho Corporation, Chennai, India

### Creator Onpremise - Multi-tenant Low Code Platform

- Built **OnPremise version of Zoho Creator from scratch**, establishing foundation for enterprise deployment
- Developed a framework to upgrade from old version of product to new, ensuring backward compatibility
- Integrated Authentication service for unified user experience across platform
- Designed and implemented pricing module of OnPremise version
- Worked on performance optimizations: **reduced application loading latency from 30s to 30ms**, component access latency from 18s to 900ms, and achieved **67% latency reduction and 41% memory improvement**
- Actively involved in code review process and bug fixing

### On Premise - Multiple Application Server

- Designed and developed a POC model for Multiple Application server architecture
- Integrated Redis and made application stateless by storing static variables in property file
- Worked in balancing load to Multiple Application Servers using Nginx

## OPEN SOURCE CONTRIBUTIONS

---

### Apache Polaris - Active Contributor | [Pull Shark Achievement](#)

Contributing to Apache Polaris, an open-source catalog for Apache Iceberg™. Working with Dropwizard and Quarkus frameworks to build scalable catalog services. Recognized with the Pull Shark achievement for multiple merged pull requests.

## CERTIFICATIONS

---

### [Mathematics for Machine Learning](#) - Coursera

### [Java Multithreading, Concurrency & Performance Optimization](#) - Udemy

### [Docker for Developers](#) - Eduative

### [Building Scalable Data Pipelines with Kafka](#) - Eduative

## EDUCATION

---

### BACHELOR OF ENGINEERING - COMPUTER SCIENCE AND ENGINEERING

Madras Institute of Technology, Chennai, India

2014 - 2018

Relevant Coursework: Data Structures, Analysis of Algorithms, Database Management Systems, Object Oriented Analysis and Design, Operating Systems, Computer Networks